

Understanding & Evaluating WebAssembly Technologies For Serverless

M Ali Ashraf Mohammad
Lawrence Wang


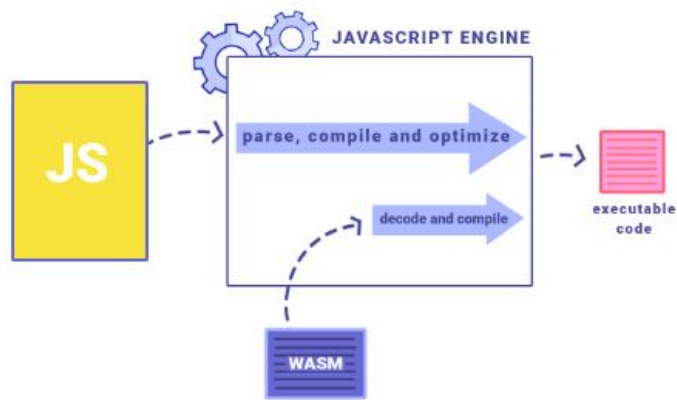


What is WebAssembly?

“WebAssembly (abbreviated Wasm) is a binary instruction format for a stack-based virtual machine.”

Wasm is designed as a portable compilation target for programming languages, enabling deployment on the web for client and server applications.

Why?

- What was it trying to solve?
- How it did it? - IR
- Being able to code in “any” programming language, compile it and run on the browsers.



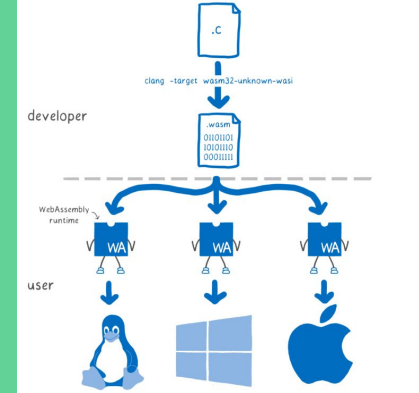
Language support to wasm & usage


- WebAssembly runs in all major browsers and in all platforms. Developers can reasonably assume WebAssembly support is anywhere JavaScript is available.
- With around 40 languages that can compile to WebAssembly, developers can finally use their favorite language on the Web.

WebAssembly has been successfully deployed in the real world, too:

- eBay implemented a [universal barcode scanner](#).
- [Google Earth](#) can run in any browser now, thanks to WebAssembly.
- The Doom 3 engine has also been ported to WebAssembly. You can play the [demo](#) online.
- Autodesk ported [AutoCad](#) to web browsers using WebAssembly.

What is Wasm being used for now? - beyond the browser

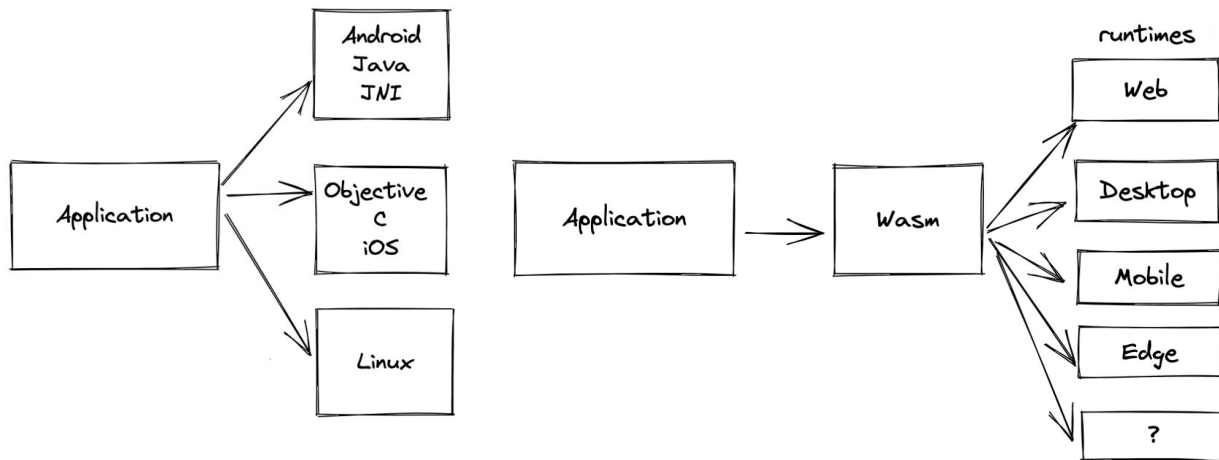


 **Solomon Hykes** / @shykes@hachyderm.io
@solomonstre

If WASM+WASI existed in 2008, we wouldn't have needed to create Docker. That's how important it is. Webassembly on the server is the future of computing. A standardized system interface was the missing link. Let's hope WASI is up to the task!

Wasm future

- Compile once, run anywhere
- Mix and match “component” tools of any language
- Secure monolithic applications with plug-ins
- Serverless & FaaS



Wasm for serverless

- Current solution: Docker containers
 - Isolation overhead - ~100s milliseconds latency for containers
 - Relatively large memory footprint
 - Inefficient state sharing between containers
- WebAssembly
 - Safety guarantee via software fault isolation
 - Memory isolation with per process contiguous memory block allocation
- Problems to solve:
 - Lessen memory restrictions for state sharing
 - Provide interface for OS operations
 - Standardized orchestration


What is our objective?

- Understand the tradeoffs between different Wasm runtimes
- Evaluate performance metrics of platforms
 - Cold start time
 - Memory consumption
 - Network usage
 - Requests per second
- Verify the accuracy of Wasm performance advertisements
- Compare Wasm options with existing serverless technologies*
 - Knative
 - SPRIGHT
- Goal: build infrastructure around runtimes to extract metrics

*Broader objective of Abhishek Sharma's project

Flavours

- [wasmer](#)
- [wasmtime](#)
- [lucet](#)
- [WAVM](#)
- [Second State VM](#)
- [V8 \(nodeJS\)](#)
- [dapr-wasm](#)
- [Intel WAMR](#)
- [wasm3](#)



What are we evaluating?

Metrics

Platforms

- ❑ CPU Utilisation, Cycles
 - ❑ Memory Footprint
 - ❑ Start-up times
-
- ❑ Wasm Flavours (Sledge, Faasm, Wasmtime*).
 - ❑ Existing serverless technologies (SPRIGHT & Knative).
-

Related Works

An Execution Model for Serverless Functions at the Edge

Adam Hall
Georgia Institute of Technology
Atlanta, Georgia
ach@gatech.edu

Umakishore Ramachandran
Georgia Institute of Technology
Atlanta, Georgia
rama@gatech.edu

A lightweight design for serverless Function-as-a- Service

Ju Long, Texas State University, Hung-Ying Tai, Shen-Ta Hsieh, and Michael Juntao Yuan, Second State LLC

Leaps and bounds: Analyzing WebAssembly's performance with a focus on bounds checking

Raven Szewczyk
University of Edinburgh
Edinburgh, UK
Raven.Szewczyk@ed.ac.uk

Antonio Barbalace
University of Edinburgh
Edinburgh, UK
Antonio.Barbalace@ed.ac.uk

Kim Stonehouse
University of Edinburgh
Edinburgh, UK
Kim.Stonehouse@ed.ac.uk

Tom Spink
University of St Andrews
St Andrews, UK
tcs6@st-andrews.ac.uk